Analysis of Redditor Reliability

Kashev Dalmia[1], Ryan Freedman[1], and Terence Nip[1],
{dalmia3, rtfreed2, nip2}@illinois.edu

[1]University of Illinois at Urbana-Champaign

ABSTRACT
In this paper, we present a system by which to evaluate the reliability of the users of the popular social network reddit. In the past, reddit has had numerous identity issues. However, through the efforts of both reddit and the userbase itself, it is becoming a place where users come to read and discuss news. Thus, there is a growing need to evaluate the reliability of the suppliers of information on reddit. We first collect features of both reliable and unreliable users based on their contributions, and most importantly, the reaction of the community to their contributions. We then use machine learning techniques to train a regression model to assign a reliability score to an arbitrary reddit user.

1. INTRODUCTION
Lorem ipsum dolor sit amet, consectetuer adipiscing elit. Ut purus elit, vestibulum ut, placerat ac, adipiscing vitae, felis. Curabitur dictum gravida mauris. Nam arcu libero, nonummy eget, consectetuer id, vulputate a, magna. Donec vehicula augue eu neque. Pellentesque habitant morbi tristique senectus et netus et malesuada fames ac turpis egestas. Mauris ut leo. Cras viverra metus rhoncus sem. Nulla et lectus vestibulum urna fringilla ultrices. Phasellus eu tellus sit amet tortor gravida placerat. Integer sapien est, iaculis in, pretium quis, viverra ac, nunc. Praesent eget sem vel leo ultrices bibendum. Aenean faucibus. Morbi dolor nulla, malesuada eu, pulvinar at, mollis ac, nulla. Curabitur auctor semper nulla. Donec varius orci eget risus. Duis nibh mi, congue eu, accumsan eleifend, sagittis quis, diam. Duis eget orci sit amet orci dignissim rutrum.

Something something boston bombing [3].

Something something news doesn't always go to the top [1]

2. RELATED WORK
Lorem ipsum dolor sit amet, consectetuer adipiscing elit. Ut purus elit, vestibulum ut, placerat ac, adipiscing vitae, felis. Curabitur dictum gravida mauris. Nam arcu libero, nonummy eget, consectetuer id, vulputate a, magna. Donec vehicula augue eu neque. Pellentesque habitant morbi tristique senectus et netus et malesuada fames ac turpis egestas. Mauris ut leo. Cras viverra metus rhoncus sem. Nulla et lectus vestibulum urna fringilla ultrices. Phasellus eu tellus sit amet tortor gravida placerat. Integer sapien est, iaculis in, pretium quis, viverra ac, nunc. Praesent eget sem vel leo ultrices bibendum. Aenean faucibus. Morbi dolor nulla, malesuada eu, pulvinar at, mollis ac, nulla. Curabitur auctor semper nulla. Donec varius orci eget risus. Duis nibh mi, congue eu, accumsan eleifend, sagittis quis, diam. Duis eget orci sit amet orci dignissim rutrum.

Something something only a single subreddit [2]

3. DESIGN
In this section, we show the design of our system. Additionally, we discuss some of the challenges, limitations, and design decisions that went into making it. Finally, we discuss some of the details of the implementation.

3.1 Gathering reddit Usernames
The first limitation of the reddit API is that usernames are an 'open secret'. If one has the user name of an account, public information about that account can be retrieved, but there is no way to directly get usernames. Instead, what we were forced to do was scrape the posts of popular subreddits and get the usernames of the author of every post and comment. In doing so, we were able to collect over 150K usernames. Of the usernames we collected, we randomly selected around 2K to fully gather data on and run our regression model on.

One issue with this approach is that this makes it impossible to identify non- participants. If a user never comments on a post, or posts a post themselves, there is no way to know that that user exists. This is unfortunately an insurmountable limitation. Instead, we chose only to find good and bad users to train our classifier, and ignore non-participants.

3.2 reddit API Limits
reddit has an API limit of 30 requests per minute. We discovered this limit is not strictly enforced, but in order to be good citizens and as to not get our access revoked, we knew we had to design around this constraint. In order to speed up our ability to access user data (as well as change the features that we used; see Section 3.4), we crawled user

data and put the raw, unmodified data into a MongoDB instance. This MongoDB served as a cache for the system. Not were we able to store raw data from reddit API calls, we were also able to cache results from more computationally intensive features.

3.3 Establishing A Ground Truth

Establishing ground truth was done in two stages. The first stage was to find reliable users. Finding these users was trivial as the site rewards positive behavior through karma, which leads to increased visibility. From there the users could be filtered by their level of contribution manually. Moderators from various communities were also taken for their work in helping the community. These users were used in our training set for a +1 reliability score.

The second stage was to find users who were unreliable or detrimental to the community. Eventually, we discovered subreddits dedicated to weeding out users that didn't contribute and various posts that detailed accounts that were used to abuse the community. These were used as our training set for a -1 reliability score.

3.4 Picking User Features: Exploratory Data Analysis

Before having a trained regression model, there was no way to tell what characteristics of reddit users would be important. However, one can't get this regression model without these features. So, we decided to pick a wide variety of features generated from the raw data, and figured out what was important later by looking at our results. Here, we present features which we suspected would be important. A discussion of what features were actually important is in **??**.

3.4.1 reddit Karma and Derivatives

TODO

3.4.2 reddit Gold and Gilded Content

TODO

3.4.3 Natural Language Processing

TODO

3.5 Picking a Regression Model

Something something neural nets give no insight

something something decision trees

something something random forest

something something darkside

4. EVALUATION AND RESULTS

From our trained random forest regression model, we get a picture of the way that redditors are. We get a glimpse of what useful, contributing redditors look like, and what bad, non-contributing redditors look like.

Table 1: The features used to create the regression model

Feature Description	Importance %
Link Karma	53.13
Comment Karma	18.41
Average Karma per Post	13.34
Number of Total Posts	9.45
% of Comment Karma - Top 100 Subreddits	2.08
Has Verified Email	0.93
Average Comment Karma per Comment	0.75
% of Comment Karma - Top 10 Subreddits	0.33
Unique Words / Total Number Words	0.29
Number of Total Comments	0.20
Time Account Created	0.17
% of Comment Karma - Trusted Subreddits	0.15
Flesch–Kincaid Readability of Comments	0.15
Is Reddit Gold	0.14
% of Post Karma - Top 50 Subreddits	0.10
% of Post Karma - Top 100 Subreddits	0.10
% of Comment Karma - Top 50 Subreddits	0.08
% of Post Karma - Trusted Subreddits	0.07
% of Comment Karma - Top 25 Subreddits	0.06
% of Swear Words Used in Comments	0.03
% of Post Karma - Top 25 Subreddits	0.03
% of Post Karma - Top 10 Subreddits	0.02
Number of Gilded Posts	0.00
% of Posts Gilded	0.00
Number of Gilded Comments	0.00
% of Comments Gilded	0.00

We collected data on around two thousand redditors, and ran their data through our regressive model to get a reliability score $-1 \leq s_r \leq 1$. Then, we re-correlate this score with input features to intuitively see what features are important or not, and what features indicated useful and not-useful redditors.

TODO

5. CONCLUSION

Lorem ipsum dolor sit amet, consectetuer adipiscing elit. Ut purus elit, vestibulum ut, placerat ac, adipiscing vitae, felis. Curabitur dictum gravida mauris. Nam arcu libero, nonummy eget, consectetuer id, vulputate a, magna. Donec vehicula augue eu neque. Pellentesque habitant morbi tristique senectus et netus et malesuada fames ac turpis egestas. Mauris ut leo. Cras viverra metus rhoncus sem. Nulla et lectus vestibulum urna fringilla ultrices. Phasellus eu tellus sit amet tortor gravida placerat. Integer sapien est, iaculis in, pretium quis, viverra ac, nunc. Praesent eget sem vel leo ultrices bibendum. Aenean faucibus. Morbi dolor nulla, malesuada eu, pulvinar at, mollis ac, nulla. Curabitur auctor semper nulla. Donec varius orci eget risus. Duis nibh mi, congue eu, accumsan eleifend, sagittis quis, diam. Duis eget orci sit amet orci dignissim rutrum.

5.1 Future Work

Something something number of posts

Something something many weaknesses

Given the set of users we have obtained in conjunction with the variability that exists among the various subsets of reddit users there exists weaknesses in our analysis. However, given the current state of research of reddit and the increasing reliance of individuals on reddit as a source of reliable in-

Figure 1: The distribution of the reliability score s_r of the sampled redditors, binned into two bins.


Figure 2: The distribution of the reliability score s_r of the sampled redditors, binned into twenty bins.

Figure 3: The reliability score s_r plotted against the percentage of gilded posts.

Figure 4: The reliability score s_r plotted against the number posts the redditor has made.

formation our contribution noteworthy for the purpose of exploration with respect to reddit as an up-and-coming source of information.

Something something exploratory/first exploration of field

TODO

5.2 Open Issues

Something something hard to establish ground truth

Voting on reddit will always be somewhat ambiguous

6. ACKNOWLEDGMENTS

7. REFERENCES

[1] E. Gilbert. Widespread underprovision on reddit. In *Proceedings of the 2013 Conference on Computer Supported Cooperative Work*, CSCW '13, pages 803–808, New York, NY, USA, 2013. ACM.

[2] A. Leavitt and J. A. Clark. Upvoting hurricane sandy: Event-based news production processes on a social news site. In *Proceedings of the SIGCHI Conference on Human Factors in Computing Systems*, CHI '14, pages 1495–1504, New York, NY, USA, 2014. ACM.

[3] L. Potts and A. Harrison. Interfaces as rhetorical constructions: Reddit and 4chan during the boston marathon bombings. In *Proceedings of the 31st ACM International Conference on Design of Communication*, SIGDOC '13, pages 143–150, New York, NY, USA, 2013. ACM.

Figure 5: The reliability score s_r plotted against the number posts the redditor has made.


Figure 6: The reliability score s_r plotted against the average Karma per post the redditor has made.


Figure 8: The distribution of the reliability score s_r of the sampled redditors, binned into twenty bins, separated by if they have a verified email address or not.


Figure 7: The reliability score s_r plotted against the Flesch–Kincaid readability of their comments.